

```

1  /*
   * linux/mm/memory.c
   *
   * Copyright (C) 1991, 1992, 1993, 1994 Linus Torvalds
5  */

/*
 * demand-loading started 01.12.91 – seems it is high on the list of
 * things wanted, and it should be easy to implement. – Linus
10 */

/*
 * Ok, demand-loading was easy, shared pages a little bit trickier. Shared
 * pages started 02.12.91, seems to work. – Linus.
15 */
 * Tested sharing by executing about 30 /bin/sh: under the old kernel it
 * would have taken more than the 6M I have free, but it worked well as
 * far as I could see.
 *
20 * Also corrected some "invalidate()"s – I wasn't doing enough of them.
 */

/*
 * Real VM (paging to/from disk) started 18.12.91. Much more work and
25 * thought has to go into this. Oh, well..
 * 19.12.91 – works, somewhat. Sometimes I get faults, don't know why.
 * Found it. Everything seems to work now.
 * 20.12.91 – Ok, making the swap-device changeable like the root.
 */

30 /*
 * 05.04.94 – Multi-page memory management added for v1.1.
 * Idea by Alex Bligh (alex@cconcepts.co.uk)
 *
35 * 16.07.99 – Support of BIGMEM added by Gerhard Wichert, Siemens AG
 * (Gerhard.Wichert@pdb.siemens.de)
 *
 * Aug/Sep 2004 Changed to four level page tables (Andi Kleen)
 */
40
#include <linux/kernel_stat.h>
#include <linux/mm.h>
#include <linux/hugetlb.h>
#include <linux/mman.h>
45 #include <linux/swap.h>
#include <linux/highmem.h>
#include <linux/pagemap.h>
#include <linux/rmap.h>
#include <linux/acct.h>
50 #include <linux/module.h>
#include <linux/init.h>

#include <asm/pgalloc.h>
#include <asm/uaccess.h>
55 #include <asm/tlb.h>
#include <asm/tlbflush.h>
#include <asm/pgtable.h>

#include <linux/swapops.h>
60 #include <linux/elf.h>

#ifdef CONFIG_DISCONTIGMEM
/* use the per-pgdat data instead for discontigmem – mbligh */

```

```

unsigned long max_mapnr;
65 struct page *mem_map;

EXPORT_SYMBOL(max_mapnr);
EXPORT_SYMBOL(mem_map);
#endif

70 unsigned long num_physpages;
/*
 * A number of key systems in x86 including ioremap() rely on the assumption
 * that high_memory defines the upper bound on direct map memory, then end
75 * of ZONE_NORMAL. Under CONFIG_DISCONTIG this means that max_low_pfn and
 * highstart_pfn must be the same; there must be no gap between ZONE_NORMAL
 * and ZONE_HIGHMEM.
 */
void * high_memory;
80 unsigned long vmalloc_earlyreserve;

EXPORT_SYMBOL(num_physpages);
EXPORT_SYMBOL(high_memory);
EXPORT_SYMBOL(vmalloc_earlyreserve);

85 /*
 * Note: this doesn't free the actual pages themselves. That
 * has been handled earlier when unmapping all the memory regions.
 */
90 static inline void clear_pmd_range(struct mmu_gather *tlb, pmd_t *pmd, unsigned long start, unsigned long end)
{
    struct page *page;

    if (pmd_none(*pmd))
95         return;
    if (unlikely(pmd_bad(*pmd))) {
        pmd_ERROR(*pmd);
        pmd_clear(pmd);
        return;
100    }
    if (!((start | end) & ~PMD_MASK)) {
        /* Only clear full, aligned ranges */
        page = pmd_page(*pmd);
        pmd_clear(pmd);
105        dec_page_state(nr_page_table_pages);
        tlb->mm->nr_ptes--;
        pte_free_tlb(tlb, page);
    }
}

110 static inline void clear_pud_range(struct mmu_gather *tlb, pud_t *pud, unsigned long start, unsigned long end)
{
    unsigned long addr = start, next;
    pmd_t *pmd, *__pmd;

115    if (pud_none(*pud))
        return;
    if (unlikely(pud_bad(*pud))) {
        pud_ERROR(*pud);
        pud_clear(pud);
120        return;
    }

    pmd = __pmd = pmd_offset(pud, start);
125    do {
        next = (addr + PMD_SIZE) & PMD_MASK;

```

```

        if (next > end || next <= addr)
            next = end;

130         clear_pmd_range(tlb, pmd, addr, next);
        pmd++;
        addr = next;
    } while (addr && (addr < end));

135     if (!(start | end) & ~PUD_MASK) {
        /* Only clear full, aligned ranges */
        pud_clear(pud);
        pmd_free_tlb(tlb, __pmd);
    }

140 }

static inline void clear_pgd_range(struct mmu_gather *tlb, pgd_t *pgd, unsigned long start, unsigned long end)
{
145     unsigned long addr = start, next;
    pud_t *pud, *__pud;

    if (pgd_none(*pgd))
        return;
150     if (unlikely(pgd_bad(*pgd))) {
        pgd_ERROR(*pgd);
        pgd_clear(pgd);
        return;
    }

155     pud = __pud = pud_offset(pgd, start);
    do {
        next = (addr + PUD_SIZE) & PUD_MASK;
        if (next > end || next <= addr)
160             next = end;

        clear_pud_range(tlb, pud, addr, next);
        pud++;
        addr = next;
165     } while (addr && (addr < end));

    if (!(start | end) & ~PGDIR_MASK) {
        /* Only clear full, aligned ranges */
        pgd_clear(pgd);
        pud_free_tlb(tlb, __pud);
170     }
}

/*
175  * This function clears user-level page tables of a process.
  *
  * Must be called with pagetable lock held.
  */
void clear_page_range(struct mmu_gather *tlb, unsigned long start, unsigned long end)
180 {
    unsigned long addr = start, next;
    pgd_t *pgd = pgd_offset(tlb->mm, start);
    unsigned long i;

185     for (i = pgd_index(start); i <= pgd_index(end-1); i++) {
        next = (addr + PGDIR_SIZE) & PGDIR_MASK;
        if (next > end || next <= addr)
            next = end;
    }
}

```

```

190         clear_pgd_range(tlb, pgd, addr, next);
        pgd++;
        addr = next;
    }
}

195 pte_t fastcall * pte_alloc_map(struct mm_struct *mm, pmd_t *pmd, unsigned long address)
{
    if (!pmd_present(*pmd)) {
        struct page *new;

200         spin_unlock(&mm->page_table_lock);
        new = pte_alloc_one(mm, address);
        spin_lock(&mm->page_table_lock);
        if (!new)
205             return NULL;

        /*
         * Because we dropped the lock, we should re-check the
         * entry, as somebody else could have populated it..
         */
210         if (pmd_present(*pmd)) {
            pte_free(new);
            goto out;
        }
        mm->nr_ptes++;
        inc_page_state(nr_page_table_pages);
        pmd_populate(mm, pmd, new);
    }

    out:
        return pte_offset_map(pmd, address);
220 }

pte_t fastcall * pte_alloc_kernel(struct mm_struct *mm, pmd_t *pmd, unsigned long address)
{
    if (!pmd_present(*pmd)) {
225         pte_t *new;

        spin_unlock(&mm->page_table_lock);
        new = pte_alloc_one_kernel(mm, address);
        spin_lock(&mm->page_table_lock);
        if (!new)
230             return NULL;

        /*
         * Because we dropped the lock, we should re-check the
         * entry, as somebody else could have populated it..
         */
235         if (pmd_present(*pmd)) {
            pte_free_kernel(new);
            goto out;
        }
        pmd_populate_kernel(mm, pmd, new);
    }

    out:
        return pte_offset_kernel(pmd, address);
245 }

/*
 * copy one vm_area from one task to the other. Assumes the page tables
 * already present in the new task to be cleared in the whole range
250 * covered by this vma.
 *
 * dst->page_table_lock is held on entry and exit,

```

```

    * but may be dropped within p[mg]d_alloc() and pte_alloc_map().
    */
255 static inline void
copy_swap_pte(struct mm_struct *dst_mm, struct mm_struct *src_mm, pte_t pte)
{
    if (pte_file(pte))
260         return;
    swap_duplicate(pte_to_swp_entry(pte));
    if (list_empty(&dst_mm->mmlist)) {
        spin_lock(&mmlist_lock);
        list_add(&dst_mm->mmlist, &src_mm->mmlist);
265         spin_unlock(&mmlist_lock);
    }
}

static inline void
270 copy_one_pte(struct mm_struct *dst_mm, struct mm_struct *src_mm,
               pte_t *dst_pte, pte_t *src_pte, unsigned long vm_flags,
               unsigned long addr)
{
    pte_t pte = *src_pte;
275     struct page *page;
    unsigned long pfn;

    /* pte contains position in swap, so copy. */
    if (!pte_present(pte)) {
280         copy_swap_pte(dst_mm, src_mm, pte);
        set_pte(dst_pte, pte);
        return;
    }
    pfn = pte_pfn(pte);
285     /* the pte points outside of valid memory, the
     * mapping is assumed to be good, meaningful
     * and not mapped via rmap - duplicate the
     * mapping as is.
     */
    page = NULL;
290     if (pfn_valid(pfn))
        page = pfn_to_page(pfn);

    if (!page || PageReserved(page)) {
295         set_pte(dst_pte, pte);
        return;
    }

    /*
300     * If it's a COW mapping, write protect it both
     * in the parent and the child
     */
    if ((vm_flags & (VM_SHARED | VM_MAYWRITE)) == VM_MAYWRITE) {
305         ptep_set_wrprotect(src_pte);
        pte = *src_pte;
    }

    /*
310     * If it's a shared mapping, mark it clean in
     * the child
     */
    if (vm_flags & VM_SHARED)
        pte = pte_mkclean(pte);
    pte = pte_mkold(pte);
315     get_page(page);

```

```

dst_mm->rss++;
if (PageAnon(page))
    dst_mm->anon_rss++;
set_pte(dst_pte, pte);
page_dup_rmap(page);
320 }

static int copy_pte_range(struct mm_struct *dst_mm, struct mm_struct *src_mm,
    pmd_t *dst_pmd, pmd_t *src_pmd, struct vm_area_struct *vma,
    unsigned long addr, unsigned long end)
325 {
    pte_t *src_pte, *dst_pte;
    pte_t *s, *d;
    unsigned long vm_flags = vma->vm_flags;

330
    d = dst_pte = pte_alloc_map(dst_mm, dst_pmd, addr);
    if (!dst_pte)
        return -ENOMEM;

335
    spin_lock(&src_mm->page_table_lock);
    s = src_pte = pte_offset_map_nested(src_pmd, addr);
    for (; addr < end; addr += PAGE_SIZE, s++, d++) {
        if (pte_none(*s))
            continue;
340
        copy_one_pte(dst_mm, src_mm, d, s, vm_flags, addr);
    }
    pte_unmap_nested(src_pte);
    pte_unmap(dst_pte);
    spin_unlock(&src_mm->page_table_lock);
345
    cond_resched_lock(&dst_mm->page_table_lock);
    return 0;
}

static int copy_pmd_range(struct mm_struct *dst_mm, struct mm_struct *src_mm,
    pud_t *dst_pud, pud_t *src_pud, struct vm_area_struct *vma,
    unsigned long addr, unsigned long end)
350 {
    pmd_t *src_pmd, *dst_pmd;
    int err = 0;
    unsigned long next;

355
    src_pmd = pmd_offset(src_pud, addr);
    dst_pmd = pmd_alloc(dst_mm, dst_pud, addr);
    if (!dst_pmd)
        return -ENOMEM;

360
    for (; addr < end; addr = next, src_pmd++, dst_pmd++) {
        next = (addr + PMD_SIZE) & PMD_MASK;
        if (next > end || next <= addr)
            next = end;
365
        if (pmd_none(*src_pmd))
            continue;
        if (pmd_bad(*src_pmd)) {
            pmd_ERROR(*src_pmd);
            pmd_clear(src_pmd);
            continue;
370
        }
        err = copy_pte_range(dst_mm, src_mm, dst_pmd, src_pmd,
                                                                    vma, addr, next);
375
        if (err)
            break;
    }
    return err;
}

```

```

    }
380
    static int copy_pud_range(struct mm_struct *dst_mm, struct mm_struct *src_mm,
                             pgd_t *dst_pgd, pgd_t *src_pgd, struct vm_area_struct *vma,
                             unsigned long addr, unsigned long end)
    {
385
        pud_t *src_pud, *dst_pud;
        int err = 0;
        unsigned long next;

        src_pud = pud_offset(src_pgd, addr);
390
        dst_pud = pud_alloc(dst_mm, dst_pgd, addr);
        if (!dst_pud)
            return -ENOMEM;

        for (; addr < end; addr = next, src_pud++, dst_pud++) {
395
            next = (addr + PUD_SIZE) & PUD_MASK;
            if (next > end || next <= addr)
                next = end;
            if (pud_none(*src_pud))
                continue;
400
            if (pud_bad(*src_pud)) {
                pud_ERROR(*src_pud);
                pud_clear(src_pud);
                continue;
            }
405
            err = copy_pmd_range(dst_mm, src_mm, dst_pud, src_pud,
                                vma, addr, next);

            if (err)
                break;
        }
410
        return err;
    }

    int copy_page_range(struct mm_struct *dst, struct mm_struct *src,
                        struct vm_area_struct *vma)
415 {
        pgd_t *src_pgd, *dst_pgd;
        unsigned long addr, start, end, next;
        int err = 0;

420
        if (is_vm_hugetlb_page(vma))
            return copy_hugetlb_page_range(dst, src, vma);

        start = vma->vm_start;
        src_pgd = pgd_offset(src, start);
425
        dst_pgd = pgd_offset(dst, start);

        end = vma->vm_end;
        addr = start;
        while (addr && (addr < end-1)) {
430
            next = (addr + PGDIR_SIZE) & PGDIR_MASK;
            if (next > end || next <= addr)
                next = end;
            if (pgd_none(*src_pgd))
                goto next_pgd;
435
            if (pgd_bad(*src_pgd)) {
                pgd_ERROR(*src_pgd);
                pgd_clear(src_pgd);
                goto next_pgd;
            }
440
            err = copy_pud_range(dst, src, dst_pgd, src_pgd,
                                vma, addr, next);

```

```

        if (err)
            break;

445 next_pgd:
        src_pgd++;
        dst_pgd++;
        addr = next;
    }

450
    return err;
}

static void zap_pte_range(struct mmu_gather *tlb,
455      pmd_t *pmd, unsigned long address,
      unsigned long size, struct zap_details *details)
{
    unsigned long offset;
    pte_t *ptep;

460
    if (pmd_none(*pmd))
        return;
    if (unlikely(pmd_bad(*pmd))) {
        pmd_ERROR(*pmd);
        pmd_clear(pmd);
465        return;
    }
    ptep = pte_offset_map(pmd, address);
    offset = address & ~PMD_MASK;
470    if (offset + size > PMD_SIZE)
        size = PMD_SIZE - offset;
    size &= PAGE_MASK;
    if (details && !details->check_mapping && !details->nonlinear_vma)
        details = NULL;
475    for (offset=0; offset < size; ptep++, offset += PAGE_SIZE) {
        pte_t pte = *ptep;
        if (pte_none(pte))
            continue;
        if (pte_present(pte)) {
480            struct page *page = NULL;
            unsigned long pfn = pte_pfn(pte);
            if (pfn_valid(pfn)) {
                page = pfn_to_page(pfn);
                if (PageReserved(page))
485                    page = NULL;
            }
            if (unlikely(details) && page) {
                /*
                 * unmap_shared_mapping_pages() wants to
490                 * invalidate cache without truncating:
                 * unmap shared but keep private pages.
                 */
                if (details->check_mapping &&
                    details->check_mapping != page->mapping)
495                    continue;
                /*
                 * Each page->index must be checked when
                 * invalidating or truncating nonlinear.
                 */
500                if (details->nonlinear_vma &&
                    (page->index < details->first_index ||
                     page->index > details->last_index))
                    continue;
            }
        }
    }
}

```



```

505         pte = ptep_get_and_clear(ptep);
        tlb_remove_tlb_entry(tlb, ptep, address+offset);
        if (unlikely(!page))
            continue;
        if (unlikely(details) && details->nonlinear_vma
510             && linear_page_index(details->nonlinear_vma,
                address+offset) != page->index)
            set_pte(ptep, pgoff_to_pte(page->index));
        if (pte_dirty(pte))
            set_page_dirty(page);
515        if (PageAnon(page))
            tlb->mm->anon_rss--;
        else if (pte_young(pte))
            mark_page_accessed(page);
        tlb->freed++;
520        page_remove_rmap(page);
        tlb_remove_page(tlb, page);
        continue;
    }
    /*
525     * If details->check_mapping, we leave swap entries;
     * if details->nonlinear_vma, we leave file entries.
     */
    if (unlikely(details))
        continue;
530    if (!pte_file(pte))
        free_swap_and_cache(pte_to_swp_entry(pte));
    pte_clear(ptep);
}
pte_unmap(ptep-1);
535 }

static void zap_pmd_range(struct mmu_gather *tlb,
                        pud_t *pud, unsigned long address,
                        unsigned long size, struct zap_details *details)
540 {
    pmd_t * pmd;
    unsigned long end;

    if (pud_none(*pud))
545        return;
    if (unlikely(pud_bad(*pud))) {
        pud_ERROR(*pud);
        pud_clear(pud);
        return;
550    }
    pmd = pmd_offset(pud, address);
    end = address + size;
    if (end > ((address + PUD_SIZE) & PUD_MASK))
        end = ((address + PUD_SIZE) & PUD_MASK);
555    do {
        zap_pte_range(tlb, pmd, address, end - address, details);
        address = (address + PMD_SIZE) & PMD_MASK;
        pmd++;
    } while (address && (address < end));
560 }

static void zap_pud_range(struct mmu_gather *tlb,
                        pgd_t *pgd, unsigned long address,
                        unsigned long end, struct zap_details *details)
565 {
    pud_t * pud;

```

```

    if (pgd_none(*pgd))
        return;
570 if (unlikely(pgd_bad(*pgd))) {
        pgd_ERROR(*pgd);
        pgd_clear(pgd);
        return;
    }
575 pud = pud_offset(pgd, address);
    do {
        zap_pmd_range(tlb, pud, address, end - address, details);
        address = (address + PUD_SIZE) & PUD_MASK;
        pud++;
580    } while (address && (address < end));
}

static void unmap_page_range(struct mmu_gather *tlb,
                            struct vm_area_struct *vma, unsigned long address,
585                            unsigned long end, struct zap_details *details)
{
    unsigned long next;
    pgd_t *pgd;
    int i;

590    BUG_ON(address >= end);
    pgd = pgd_offset(vma->vm_mm, address);
    tlb_start_vma(tlb, vma);
    for (i = pgd_index(address); i <= pgd_index(end-1); i++) {
595        next = (address + PGDIR_SIZE) & PGDIR_MASK;
        if (next <= address || next > end)
            next = end;
        zap_pud_range(tlb, pgd, address, next, details);
        address = next;
600        pgd++;
    }
    tlb_end_vma(tlb, vma);
}

605 #ifdef CONFIG_PREEMPT
# define ZAP_BLOCK_SIZE (8 * PAGE_SIZE)
#else
/* No preempt: go for improved straight-line efficiency */
# define ZAP_BLOCK_SIZE (1024 * PAGE_SIZE)
610 #endif

/*
 * unmap_vmas - unmap a range of memory covered by a list of vma's
 * @tlbp: address of the caller's struct mmu_gather
615 * @mm: the controlling mm_struct
 * @vma: the starting vma
 * @start_addr: virtual address at which to start unmapping
 * @end_addr: virtual address at which to end unmapping
 * @nr_accounted: Place number of unmapped pages in vm-accountable vma's here
620 * @details: details of nonlinear truncation or shared cache invalidation
 *
 * Returns the number of vma's which were covered by the unmapping.
 *
 * Unmap all pages in the vma list. Called under page_table_lock.
625 *
 * We aim to not hold page_table_lock for too long (for scheduling latency
 * reasons). So zap pages in ZAP_BLOCK_SIZE bytecounts. This means we need to
 * return the ending mmu_gather to the caller.
 *
630 * Only addresses between 'start' and 'end' will be unmapped.

```

```

*
* The VMA list must be sorted in ascending virtual address order.
*
* unmap_vmas() assumes that the caller will flush the whole unmapped address
635 * range after unmap_vmas() returns. So the only responsibility here is to
* ensure that any thus-far unmapped pages are flushed before unmap_vmas()
* drops the lock and schedules.
*/
int unmap_vmas(struct mmu_gather **tlbp, struct mm_struct *mm,
640 struct vm_area_struct *vma, unsigned long start_addr,
unsigned long end_addr, unsigned long *nr_accounted,
struct zap_details *details)
{
    unsigned long zap_bytes = ZAP_BLOCK_SIZE;
645 unsigned long tlb_start = 0; /* For tlb_finish_mmu */
int tlb_start_valid = 0;
int ret = 0;
spinlock_t *i_mmap_lock = details? details->i_mmap_lock: NULL;
int fullmm = tlb_is_full_mm(*tlbp);
650
    for ( ; vma && vma->vm_start < end_addr; vma = vma->vm_next) {
        unsigned long start;
        unsigned long end;

655 start = max(vma->vm_start, start_addr);
        if (start >= vma->vm_end)
            continue;
        end = min(vma->vm_end, end_addr);
        if (end <= vma->vm_start)
660 continue;

        if (vma->vm_flags & VM_ACCOUNT)
            *nr_accounted += (end - start) >> PAGE_SHIFT;

665 ret++;
        while (start != end) {
            unsigned long block;

            if (!tlb_start_valid) {
670 tlb_start = start;
                tlb_start_valid = 1;
            }

            if (is_vm_hugetlb_page(vma)) {
675 block = end - start;
                unmap_hugepage_range(vma, start, end);
            } else {
                block = min(zap_bytes, end - start);
                unmap_page_range(*tlbp, vma, start,
680 start + block, details);
            }

            start += block;
            zap_bytes -= block;
685 if ((long)zap_bytes > 0)
                continue;

            tlb_finish_mmu(*tlbp, tlb_start, start);

690 if (need_resched() ||
                need_lockbreak(&mm->page_table_lock) ||
                (i_mmap_lock && need_lockbreak(i_mmap_lock))) {
                    if (i_mmap_lock) {

```

```

695                                     /* must reset count of rss freed */
                                     *tlbp = tlb_gather_mmu(mm, fullmm);
                                     details->break_addr = start;
                                     goto out;
                                     }
700                                 spin_unlock(&mm->page_table_lock);
                                 cond_resched();
                                 spin_lock(&mm->page_table_lock);
                                }

705                                *tlbp = tlb_gather_mmu(mm, fullmm);
                                tlb_start_valid = 0;
                                zap_bytes = ZAP_BLOCK_SIZE;
                                }
                                }

out:
710    return ret;
}

/*
 * zap_page_range - remove user pages in a given range
715 * @vma: vm_area_struct holding the applicable pages
 * @address: starting address of pages to zap
 * @size: number of bytes to zap
 * @details: details of nonlinear truncation or shared cache invalidation
 */
720 void zap_page_range(struct vm_area_struct *vma, unsigned long address,
                      unsigned long size, struct zap_details *details)
{
    struct mm_struct *mm = vma->vm_mm;
    struct mmu_gather *tlb;
725    unsigned long end = address + size;
    unsigned long nr_accounted = 0;

    if (is_vm_hugetlb_page(vma)) {
        zap_hugepage_range(vma, address, size);
730        return;
    }

    lru_add_drain();
    spin_lock(&mm->page_table_lock);
735    tlb = tlb_gather_mmu(mm, 0);
    unmap_vmas(&tlb, mm, vma, address, end, &nr_accounted, details);
    tlb_finish_mmu(tlb, address, end);
    acct_update_integrals();
    spin_unlock(&mm->page_table_lock);
740 }

/*
 * Do a quick page-table lookup for a single page.
 * mm->page_table_lock must be held.
745 */
static struct page *
__follow_page(struct mm_struct *mm, unsigned long address, int read, int write)
{
    pgd_t *pgd;
750    pud_t *pud;
    pmd_t *pmd;
    pte_t *ptep, pte;
    unsigned long pfn;
    struct page *page;

755    page = follow_huge_addr(mm, address, write);

```

```

        if (!IS_ERR(page))
            return page;

760     pgd = pgd_offset(mm, address);
    if (pgd_none(*pgd) || unlikely(pgd_bad(*pgd)))
        goto out;

    pud = pud_offset(pgd, address);
765     if (pud_none(*pud) || unlikely(pud_bad(*pud)))
        goto out;

    pmd = pmd_offset(pud, address);
    if (pmd_none(*pmd) || unlikely(pmd_bad(*pmd)))
770         goto out;
    if (pmd_huge(*pmd))
        return follow_huge_pmd(mm, address, pmd, write);

    ptep = pte_offset_map(pmd, address);
775     if (!ptep)
        goto out;

    pte = *ptep;
    pte_unmap(ptep);
780     if (pte_present(pte)) {
        if (write && !pte_write(pte))
            goto out;
        if (read && !pte_read(pte))
            goto out;
785         pfn = pte_pfn(pte);
        if (pfn_valid(pfn)) {
            page = pfn_to_page(pfn);
            if (write && !pte_dirty(pte) && !PageDirty(page))
                set_page_dirty(page);
790             mark_page_accessed(page);
            return page;
        }
    }

795 out:
    return NULL;
}

struct page *
800 follow_page(struct mm_struct *mm, unsigned long address, int write)
{
    return __follow_page(mm, address, /* read */ 0, write);
}

805 int
check_user_page_readable(struct mm_struct *mm, unsigned long address)
{
    return __follow_page(mm, address, /* read */ 1, /* write */ 0) != NULL;
}

810 EXPORT_SYMBOL(check_user_page_readable);

/*
 * Given a physical address, is there a useful struct page pointing to
815 * it? This may become more complex in the future if we start dealing
 * with IO-aperture pages for direct-IO.
 */

static inline struct page *get_page_map(struct page *page)

```

```

820 {
    if (!pfn_valid(page_to_pfn(page)))
        return NULL;
    return page;
}

825

static inline int
untouched_anonymous_page(struct mm_struct* mm, struct vm_area_struct *vma,
                          unsigned long address)

830 {
    pgd_t *pgd;
    pud_t *pud;
    pmd_t *pmd;

835
    /* Check if the vma is for an anonymous mapping. */
    if (vma->vm_ops && vma->vm_ops->nopage)
        return 0;

    /* Check if page directory entry exists. */
840
    pgd = pgd_offset(mm, address);
    if (pgd_none(*pgd) || unlikely(pgd_bad(*pgd)))
        return 1;

    pud = pud_offset(pgd, address);
845
    if (pud_none(*pud) || unlikely(pud_bad(*pud)))
        return 1;

    /* Check if page middle directory entry exists. */
    pmd = pmd_offset(pud, address);
850
    if (pmd_none(*pmd) || unlikely(pmd_bad(*pmd)))
        return 1;

    /* There is a pte slot for 'address' in 'mm'. */
    return 0;
855 }

int get_user_pages(struct task_struct *tsk, struct mm_struct *mm,
                  unsigned long start, int len, int write, int force,
                  struct page **pages, struct vm_area_struct **vmas)

860
{
    int i;
    unsigned int flags;

865
    /*
     * Require read or write permissions.
     * If 'force' is set, we only require the "MAY" flags.
     */
    flags = write ? (VM_WRITE | VM_MAYWRITE) : (VM_READ | VM_MAYREAD);
    flags &= force ? (VM_MAYREAD | VM_MAYWRITE) : (VM_READ | VM_WRITE);
870
    i = 0;

    do {
        struct vm_area_struct * vma;

875
        vma = find_extend_vma(mm, start);
        if (!vma && in_gate_area(tsk, start)) {
            unsigned long pg = start & PAGE_MASK;
            struct vm_area_struct *gate_vma = get_gate_vma(tsk);
            pgd_t *pgd;
            pud_t *pud;
            pmd_t *pmd;
880

```

```

pte_t *pte;
885 if (write) /* user gate pages are read-only */
        return i ? : -EFAULT;
if (pg > TASK_SIZE)
    pgd = pgd_offset_k(pg);
else
    pgd = pgd_offset_gate(mm, pg);
890 BUG_ON(pgd_none(*pgd));
pud = pud_offset(pgd, pg);
BUG_ON(pud_none(*pud));
pmd = pmd_offset(pud, pg);
BUG_ON(pmd_none(*pmd));
895 pte = pte_offset_map(pmd, pg);
BUG_ON(pte_none(*pte));
if (pages) {
    pages[i] = pte_page(*pte);
    get_page(pages[i]);
900 }
pte_unmap(pte);
if (vmass)
    vmass[i] = gate_vma;

i++;
905 start += PAGE_SIZE;
len--;
continue;
}

910 if (!vma || (vma->vm_flags & VM_IO)
        || !(flags & vma->vm_flags))
    return i ? : -EFAULT;

if (is_vm_hugetlb_page(vma)) {
915 i = follow_hugetlb_page(mm, vma, pages, vmass,
                        &start, &len, i);

    continue;
}
spin_lock(&mm->page_table_lock);
920 do {

    struct page *map;
    int lookup_write = write;

    cond_resched_lock(&mm->page_table_lock);
925 while (!(map = follow_page(mm, start, lookup_write))) {
        /*
         * Shortcut for anonymous pages. We don't want
         * to force the creation of pages tables for
         * insanely big anonymously mapped areas that
         * nobody touched so far. This is important
         * for doing a core dump for these mappings.
         */
        if (!lookup_write &&
            untouched_anonymous_page(mm, vma, start)) {
935 map = ZERO_PAGE(start);
            break;
        }
        spin_unlock(&mm->page_table_lock);
        switch (handle_mm_fault(mm, vma, start, write)) {
940 case VM_FAULT_MINOR:
            tsk->min_flt++;
            break;
            case VM_FAULT_MAJOR:
            tsk->maj_flt++;
945 break;

```

```

case VM_FAULT_SIGBUS:
    return i ? i : -EFAULT;
case VM_FAULT_OOM:
    return i ? i : -ENOMEM;
default:
    BUG();
}
/*
 * Now that we have performed a write fault
 * and surely no longer have a shared page we
 * shouldn't write, we shouldn't ignore an
 * unwritable page in the page table if
 * we are forcing write access.
 */
lookup_write = write && !force;
spin_lock(&mm->page_table_lock);
}
if (pages) {
    pages[i] = get_page_map(map);
    if (!pages[i]) {
        spin_unlock(&mm->page_table_lock);
        while (i--)
            page_cache_release(pages[i]);
        i = -EFAULT;
        goto out;
    }
    flush_dcache_page(pages[i]);
    if (!PageReserved(pages[i]))
        page_cache_get(pages[i]);
}
if (vmas)
    vmas[i] = vma;
i++;
start += PAGE_SIZE;
len--;
} while(len && start < vma->vm_end);
spin_unlock(&mm->page_table_lock);
} while(len);
out:
return i;
}

EXPORT_SYMBOL(get_user_pages);

990 static void zeromap_pte_range(pte_t * pte, unsigned long address,
                                unsigned long size, pgprot_t prot)
{
    unsigned long end;

    address &= ~PMD_MASK;
    end = address + size;
    if (end > PMD_SIZE)
        end = PMD_SIZE;
    do {
        pte_t zero_pte = pte_wrprotect(mk_pte(ZERO_PAGE(address), prot));
        BUG_ON(!pte_none(*pte));
        set_pte(pte, zero_pte);
        address += PAGE_SIZE;
        pte++;
    } while (address && (address < end));
}

static inline int zeromap_pmd_range(struct mm_struct *mm, pmd_t * pmd,

```



```

1010 {
    unsigned long address, unsigned long size, pgprot_t prot)

    unsigned long base, end;

    base = address & PUD_MASK;
    address &= ~PUD_MASK;
1015    end = address + size;
    if (end > PUD_SIZE)
        end = PUD_SIZE;
    do {
        pte_t * pte = pte_alloc_map(mm, pmd, base + address);
1020        if (!pte)
            return -ENOMEM;
        zeromap_pte_range(pte, base + address, end - address, prot);
        pte_unmap(pte);
        address = (address + PMD_SIZE) & PMD_MASK;
1025        pmd++;
    } while (address && (address < end));
    return 0;
}

1030 static inline int zeromap_pud_range(struct mm_struct *mm, pud_t * pud,
    unsigned long address,
    unsigned long size, pgprot_t prot)
{
    unsigned long base, end;
1035    int error = 0;

    base = address & PGDIR_MASK;
    address &= ~PGDIR_MASK;
    end = address + size;
1040    if (end > PGDIR_SIZE)
        end = PGDIR_SIZE;
    do {
        pmd_t * pmd = pmd_alloc(mm, pud, base + address);
        error = -ENOMEM;
1045        if (!pmd)
            break;
        error = zeromap_pmd_range(mm, pmd, base + address,
            end - address, prot);
        if (error)
1050            break;
        address = (address + PUD_SIZE) & PUD_MASK;
        pud++;
    } while (address && (address < end));
    return 0;
1055 }

int zeromap_page_range(struct vm_area_struct *vma, unsigned long address,
    unsigned long size, pgprot_t prot)
{
1060    int i;
    int error = 0;
    pgd_t * pgd;
    unsigned long beg = address;
    unsigned long end = address + size;
1065    unsigned long next;
    struct mm_struct *mm = vma->vm_mm;

    pgd = pgd_offset(mm, address);
    flush_cache_range(vma, beg, end);
1070    BUG_ON(address >= end);
    BUG_ON(end > vma->vm_end);

```

```

spin_lock(&mm->page_table_lock);
for (i = pgd_index(address); i <= pgd_index(end-1); i++) {
1075     pud_t *pud = pud_alloc(mm, pgd, address);
    error = -ENOMEM;
    if (!pud)
        break;
    next = (address + PGDIR_SIZE) & PGDIR_MASK;
1080     if (next <= beg || next > end)
        next = end;
    error = zeromap_pud_range(mm, pud, address,
                                next - address, prot);

    if (error)
1085         break;
    address = next;
    pgd++;
}
/*
1090  * Why flush? zeromap_pte_range has a BUG_ON for !pte_none()
  */
flush_tlb_range(vma, beg, end);
spin_unlock(&mm->page_table_lock);
return error;
1095 }

```

```

/*
 * maps a range of physical memory into the requested pages. the old
 * mappings are removed. any references to nonexistent pages results
1100 * in null mappings (currently treated as "copy-on-access")
 */

```

static inline void

remap_pte_range(pte_t * pte, **unsigned long** address, **unsigned long** size,
unsigned long pfn, pgprot_t prot)

```

1105 {
    unsigned long end;

    address &= ~PMD_MASK;
    end = address + size;
1110     if (end > PMD_SIZE)
        end = PMD_SIZE;

    do {
        BUG_ON(!pte_none(*pte));
        if (!pfn_valid(pfn) || PageReserved(pfn_to_page(pfn)))
1115             set_pte(pte, pfn_pte(pfn, prot));
        address += PAGE_SIZE;
        pfn++;
        pte++;
    } while (address && (address < end));
1120 }

```

static inline int

remap_pmd_range(**struct** mm_struct *mm, pmd_t * pmd, **unsigned long** address,
unsigned long size, **unsigned long** pfn, pgprot_t prot)

```

1125 {
    unsigned long base, end;

    base = address & PUD_MASK;
    address &= ~PUD_MASK;
1130     end = address + size;
    if (end > PUD_SIZE)
        end = PUD_SIZE;
    pfn -= (address >> PAGE_SHIFT);
    do {

```

```

1135         pte_t * pte = pte_alloc_map(mm, pmd, base + address);
        if (!pte)
            return -ENOMEM;
        remap_pte_range(pte, base + address, end - address,
                        (address >> PAGE_SHIFT) + pfn, prot);
1140     pte_unmap(pte);
        address = (address + PMD_SIZE) & PMD_MASK;
        pmd++;
    } while (address && (address < end));
    return 0;
1145 }

static inline int remap_pud_range(struct mm_struct *mm, pud_t * pud,
                                unsigned long address, unsigned long size,
                                unsigned long pfn, pgprot_t prot)
1150 {
    unsigned long base, end;
    int error;

    base = address & PGDIR_MASK;
    address &= ~PGDIR_MASK;
    end = address + size;
    if (end > PGDIR_SIZE)
        end = PGDIR_SIZE;
    pfn -= address >> PAGE_SHIFT;
1160     do {
        pmd_t *pmd = pmd_alloc(mm, pud, base+address);
        error = -ENOMEM;
        if (!pmd)
            break;
1165         error = remap_pmd_range(mm, pmd, base + address, end - address,
                                (address >> PAGE_SHIFT) + pfn, prot);
        if (error)
            break;
        address = (address + PUD_SIZE) & PUD_MASK;
        pud++;
1170     } while (address && (address < end));
    return error;
}

1175 /* Note: this is only safe if the mm semaphore is held when called. */
int remap_pfn_range(struct vm_area_struct *vma, unsigned long from,
                   unsigned long pfn, unsigned long size, pgprot_t prot)
{
    int error = 0;
    pgd_t *pgd;
    unsigned long beg = from;
    unsigned long end = from + size;
    unsigned long next;
    struct mm_struct *mm = vma->vm_mm;
1185     int i;

    pfn -= from >> PAGE_SHIFT;
    pgd = pgd_offset(mm, from);
    flush_cache_range(vma, beg, end);
1190     BUG_ON(from >= end);

    /*
     * Physically remapped pages are special. Tell the
     * rest of the world about it:
     *     VM_IO tells people not to look at these pages
     *     (accesses can have side effects).
     *     VM_RESERVED tells swapout not to try to touch

```

```

    *           this region.
    */
1200 vma->vm_flags |= VM_IO | VM_RESERVED;

spin_lock(&mm->page_table_lock);
for (i = pgd_index(beg); i <= pgd_index(end-1); i++) {
1205     pud_t *pud = pud_alloc(mm, pgd, from);
    error = -ENOMEM;
    if (!pud)
        break;
    next = (from + PGDIR_SIZE) & PGDIR_MASK;
    if (next > end || next <= from)
        next = end;
1210     error = remap_pud_range(mm, pud, from, end - from,
                             pfn + (from >> PAGE_SHIFT), prot);

    if (error)
        break;
1215     from = next;
    pgd++;
}
/*
 * Why flush? remap_pte_range has a BUG_ON for !pte_none()
 */
1220 flush_tlb_range(vma, beg, end);
spin_unlock(&mm->page_table_lock);

return error;
1225 }

```

```
EXPORT_SYMBOL(remap_pfn_range);
```

```

/*
1230 * Do pte_mkwrite, but only if the vma says VM_WRITE. We do this when
 * servicing faults for write access. In the normal case, do always want
 * pte_mkwrite. But get_user_pages can cause write faults for mappings
 * that do not have writing enabled, when used by access_process_vm.
 */

```

```

1235 static inline pte_t maybe_mkwrite(pte_t pte, struct vm_area_struct *vma)
{
    if (likely(vma->vm_flags & VM_WRITE))
        pte = pte_mkwrite(pte);
    return pte;
1240 }

```

```

/*
 * We hold the mm semaphore for reading and vma->vm_mm->page_table_lock
 */
1245 static inline void break_cow(struct vm_area_struct * vma, struct page * new_page, unsigned long address,
    pte_t * page_table)
{

```

```

    pte_t entry;

1250     flush_cache_page(vma, address);
    entry = maybe_mkwrite(pte_mkdirty(mk_pte(new_page, vma->vm_page_prot)),
                          vma);
    ptep_establish(vma, address, page_table, entry);
    update_mmu_cache(vma, address, entry);
1255 }

```

```

/*
 * This routine handles present pages, when users try to write
 * to a shared page. It is done by copying the page to a new address
1260 * and decrementing the shared-page counter for the old page.

```

```

*
* Goto-purists beware: the only reason for goto's here is that it results
* in better assembly code.. The "default" path will see no jumps at all.
*
1265 * Note that this routine assumes that the protection checks have been
* done by the caller (the low-level page fault routine in most cases).
* Thus we can safely just mark it writable once we've done any necessary
* COW.
*
1270 * We also mark the page dirty at this point even though the page will
* change only once the write actually happens. This avoids a few races,
* and potentially makes it more efficient.
*
* We hold the mm semaphore and the page_table_lock on entry and exit
1275 * with the page_table_lock released.
* /
static int do_wp_page(struct mm_struct *mm, struct vm_area_struct * vma,
                     unsigned long address, pte_t * page_table, pmd_t * pmd, pte_t pte)
{
1280     struct page *old_page, *new_page;
     unsigned long pfn = pte_pfn(pte);
     pte_t entry;

     if (unlikely(!pfn_valid(pfn))) {
1285         /*
          * This should really halt the system so it can be debugged or
          * at least the kernel stops what it's doing before it corrupts
          * data, but for the moment just pretend this is OOM.
          */
1290         pte_unmap(page_table);
         printk(KERN_ERR "do_wp_page: bogus page at address %08lx\n",
                address);
         spin_unlock(&mm->page_table_lock);
         return VM_FAULT_OOM;
1295     }
     old_page = pfn_to_page(pfn);

     if (!TestSetPageLocked(old_page)) {
         int reuse = can_share_swap_page(old_page);
1300         unlock_page(old_page);
         if (reuse) {
             flush_cache_page(vma, address);
             entry = maybe_mkwrite(pte_mkyoung(pte_mkdirty(pte)),
                                   vma);
1305             ptep_set_access_flags(vma, address, page_table, entry, 1);
             update_mmu_cache(vma, address, entry);
             pte_unmap(page_table);
             spin_unlock(&mm->page_table_lock);
             return VM_FAULT_MINOR;
1310         }
     }
     pte_unmap(page_table);

     /*
1315     * Ok, we need to copy. Oh, well..
     */
     if (!PageReserved(old_page))
         page_cache_get(old_page);
     spin_unlock(&mm->page_table_lock);
1320
     if (unlikely(anon_vma_prepare(vma)))
         goto no_new_page;
     if (old_page == ZERO_PAGE(address)) {

```

```

1325     new_page = alloc_zeroed_user_highpage(vma, address);
        if (!new_page)
            goto no_new_page;
    } else {
        new_page = alloc_page_vma(GFP_HIGHUSER, vma, address);
        if (!new_page)
1330             goto no_new_page;
        copy_user_highpage(new_page, old_page, address);
    }
    /*
     * Re-check the pte - we dropped the lock
1335     */
    spin_lock(&mm->page_table_lock);
    page_table = pte_offset_map(pmd, address);
    if (likely(pte_same(*page_table, pte))) {
        if (PageAnon(old_page))
1340             mm->anon_rss--;
        if (PageReserved(old_page)) {
            ++mm->rss;
            acct_update_integrals();
            update_mem_hiwater();
1345        } else
            page_remove_rmap(old_page);
        break_cow(vma, new_page, address, page_table);
        lru_cache_add_active(new_page);
        page_add_anon_rmap(new_page, vma, address);
1350
        /* Free the old page.. */
        new_page = old_page;
    }
    pte_unmap(page_table);
1355    page_cache_release(new_page);
    page_cache_release(old_page);
    spin_unlock(&mm->page_table_lock);
    return VM_FAULT_MINOR;

1360 no_new_page:
    page_cache_release(old_page);
    return VM_FAULT_OOM;
}

1365 /*
     * Helper functions for unmap_mapping_range().
     *
     * __ Notes on dropping i_mmap_lock to reduce latency while unmapping __
     *
1370 * We have to restart searching the prio_tree whenever we drop the lock,
     * since the iterator is only valid while the lock is held, and anyway
     * a later vma might be split and reinserted earlier while lock dropped.
     *
     * The list of nonlinear vmAs could be handled more efficiently, using
1375 * a placeholder, but handle it in the same way until a need is shown.
     * It is important to search the prio_tree before nonlinear list: a vma
     * may become nonlinear and be shifted from prio_tree to nonlinear list
     * while the lock is dropped; but never shifted from list to prio_tree.
     *
1380 * In order to make forward progress despite restarting the search,
     * vm_truncate_count is used to mark a vma as now dealt with, so we can
     * quickly skip it next time around. Since the prio_tree search only
     * shows us those vmAs affected by unmapping the range in question, we
     * can't efficiently keep all vmAs in step with mapping->truncate_count:
1385 * so instead reset them all whenever it wraps back to 0 (then go to 1).
     * mapping->truncate_count and vma->vm_truncate_count are protected by

```

```

* i_mmap_lock.
*
* In order to make forward progress despite repeatedly restarting some
1390 * large vma, note the break_addr set by unmap_vmas when it breaks out:
* and restart from that address when we reach that vma again. It might
* have been split or merged, shrunk or extended, but never shifted: so
* restart_addr remains valid so long as it remains in the vma's range.
* unmap_mapping_range forces truncate_count to leap over page-aligned
1395 * values so we can save vma's restart_addr in its truncate_count field.
*/
#define is_restart_addr(truncate_count) (!((truncate_count) & ~PAGE_MASK))

static void reset_vma_truncate_counts(struct address_space *mapping)
1400 {
    struct vm_area_struct *vma;
    struct prio_tree_iter iter;

    vma_prio_tree_foreach(vma, &iter, &mapping->i_mmap, 0, ULONG_MAX)
1405     vma->vm_truncate_count = 0;
    list_for_each_entry(vma, &mapping->i_mmap_nonlinear, shared.vm_set.list)
        vma->vm_truncate_count = 0;
}

1410 static int unmap_mapping_range_vma(struct vm_area_struct *vma,
    unsigned long start_addr, unsigned long end_addr,
    struct zap_details *details)
{
    unsigned long restart_addr;
1415    int need_break;

again:
    restart_addr = vma->vm_truncate_count;
    if (is_restart_addr(restart_addr) && start_addr < restart_addr) {
1420         start_addr = restart_addr;
        if (start_addr >= end_addr) {
            /* Top of vma has been split off since last time */
            vma->vm_truncate_count = details->truncate_count;
            return 0;
1425         }
    }

    details->break_addr = end_addr;
    zap_page_range(vma, start_addr, end_addr - start_addr, details);

1430
    /*
     * We cannot rely on the break test in unmap_vmas:
     * on the one hand, we don't want to restart our loop
     * just because that broke out for the page_table_lock;
1435     * on the other hand, it does no test when vma is small.
     */
    need_break = need_resched() ||
        need_lockbreak(details->i_mmap_lock);

1440    if (details->break_addr >= end_addr) {
        /* We have now completed this vma: mark it so */
        vma->vm_truncate_count = details->truncate_count;
        if (!need_break)
            return 0;
1445    } else {
        /* Note restart_addr in vma's truncate_count field */
        vma->vm_truncate_count = details->break_addr;
        if (!need_break)
            goto again;
    }
}

```

```

1450     }

    spin_unlock(details->i_mmap_lock);
    cond_resched();
    spin_lock(details->i_mmap_lock);
1455     return -EINTR;
}

static inline void unmap_mapping_range_tree(struct prio_tree_root *root,
                                           struct zap_details *details)

1460 {
    struct vm_area_struct *vma;
    struct prio_tree_iter iter;
    pgoff_t vba, vea, zba, zea;

1465 restart:
    vma_prio_tree_foreach(vma, &iter, root,
                          details->first_index, details->last_index) {
        /* Skip quickly over those we have already dealt with */
        if (vma->vm_truncate_count == details->truncate_count)
1470             continue;

        vba = vma->vm_pgoff;
        vea = vba + ((vma->vm_end - vma->vm_start) >> PAGE_SHIFT) - 1;
        /* Assume for now that PAGE_CACHE_SHIFT == PAGE_SHIFT */
1475        zba = details->first_index;
        if (zba < vba)
            zba = vba;
        zea = details->last_index;
        if (zea > vea)
1480             zea = vea;

        if (unmap_mapping_range_vma(vma,
                                     ((zba - vba) << PAGE_SHIFT) + vma->vm_start,
                                     ((zea - vba + 1) << PAGE_SHIFT) + vma->vm_start,
1485                                     details) < 0)
            goto restart;
    }
}

1490 static inline void unmap_mapping_range_list(struct list_head *head,
                                           struct zap_details *details)
{
    struct vm_area_struct *vma;

1495     /*
     * In nonlinear VMAs there is no correspondence between virtual address
     * offset and file offset. So we must perform an exhaustive search
     * across *all* the pages in each nonlinear VMA, not just the pages
     * whose virtual address lies outside the file truncation point.
1500     */

    restart:
    list_for_each_entry(vma, head, shared.vm_set.list) {
        /* Skip quickly over those we have already dealt with */
        if (vma->vm_truncate_count == details->truncate_count)
1505             continue;
        details->nonlinear_vma = vma;
        if (unmap_mapping_range_vma(vma, vma->vm_start,
                                     vma->vm_end, details) < 0)
            goto restart;
1510    }
}

```



```

1515 /*
1520  * unmap_mapping_range – unmap the portion of all mmaps
1525  * in the specified address_space corresponding to the specified
1530  * page range in the underlying file.
1535  * @address_space: the address space containing mmaps to be unmapped.
1540  * @holebegin: byte in first page to unmap, relative to the start of
1545  * the underlying file. This will be rounded down to a PAGE_SIZE
1550  * boundary. Note that this is different from vmtruncate(), which
1555  * must keep the partial page. In contrast, we must get rid of
1560  * partial pages.
1565  * @holelen: size of prospective hole in bytes. This will be rounded
1570  * up to a PAGE_SIZE boundary. A holelen of zero truncates to the
1575  * end of the file.
1580  * @even_cows: 1 when truncating a file, unmap even private COWed pages;
1585  * but 0 when invalidating pagecache, don't throw away private data.
1590  */
1595 void unmap_mapping_range(struct address_space *mapping,
1600 loff_t const holebegin, loff_t const holelen, int even_cows)
1605 {
1610     struct zap_details details;
1615     pgoff_t hba = holebegin >> PAGE_SHIFT;
1620     pgoff_t hlen = (holelen + PAGE_SIZE - 1) >> PAGE_SHIFT;
1625
1630     /* Check for overflow. */
1635     if (sizeof(holelen) > sizeof(hlen)) {
1640         long long holeend =
1645             (holebegin + holelen + PAGE_SIZE - 1) >> PAGE_SHIFT;
1650         if (holeend & ~(long long)ULONG_MAX)
1655             hlen = ULONG_MAX - hba + 1;
1660     }
1665
1670     details.check_mapping = even_cows? NULL: mapping;
1675     details.nonlinear_vma = NULL;
1680     details.first_index = hba;
1685     details.last_index = hba + hlen - 1;
1690     if (details.last_index < details.first_index)
1695         details.last_index = ULONG_MAX;
1700     details.i_mmap_lock = &mapping->i_mmap_lock;
1705
1710     spin_lock(&mapping->i_mmap_lock);
1715
1720     /* serialize i_size write against truncate_count write */
1725     smp_wmb();
1730     /* Protect against page faults, and endless unmapping loops */
1735     mapping->truncate_count++;
1740     /*
1745      * For archs where spin_lock has inclusive semantics like ia64
1750      * this smp_mb() will prevent to read pagetable contents
1755      * before the truncate_count increment is visible to
1760      * other cpus.
1765      */
1770     smp_mb();
1775     if (unlikely(is_restart_addr(mapping->truncate_count))) {
1780         if (mapping->truncate_count == 0)
1785             reset_vma_truncate_counts(mapping);
1790         mapping->truncate_count++;
1795     }
1800     details.truncate_count = mapping->truncate_count;
1805
1810     if (unlikely(!prio_tree_empty(&mapping->i_mmap)))
1815         unmap_mapping_range_tree(&mapping->i_mmap, &details);
1820     if (unlikely(!list_empty(&mapping->i_mmap_nonlinear)))
1825         unmap_mapping_range_list(&mapping->i_mmap_nonlinear, &details);
1830

```

```

        spin_unlock(&mapping->i_mmap_lock);
    }
    EXPORT_SYMBOL(unmap_mapping_range);

1580 /*
    * Handle all mappings that got truncated by a "truncate()"
    * system call.
    *
    * NOTE! We have to be ready to update the memory sharing
1585 * between the file and the memory map for a potential last
    * incomplete page. Ugly, but necessary.
    */
    int vmtruncate(struct inode * inode, loff_t offset)
    {
1590         struct address_space * mapping = inode->i_mapping;
        unsigned long limit;

        if (inode->i_size < offset)
            goto do_expand;

1595         /*
            * truncation of in-use swapfiles is disallowed – it would cause
            * subsequent swapout to scribble on the now-freed blocks.
            */
        if (IS_SWAPFILE(inode))
            goto out_busy;
        i_size_write(inode, offset);
        unmap_mapping_range(mapping, offset + PAGE_SIZE - 1, 0, 1);
        truncate_inode_pages(mapping, offset);
        goto out_truncate;

1605 do_expand:
        limit = current->signal->rlim[RLIMIT_FSIZE].rlim_cur;
        if (limit != RLIM_INFINITY && offset > limit)
            goto out_sig;
1610         if (offset > inode->i_sb->s_maxbytes)
            goto out_big;
        i_size_write(inode, offset);

        out_truncate:
1615         if (inode->i_op && inode->i_op->truncate)
            inode->i_op->truncate(inode);
        return 0;

        out_sig:
            send_sig(SIGXFSZ, current, 0);
1620 out_big:
            return -EFBIG;

        out_busy:
            return -ETXTBSY;
    }

1625 EXPORT_SYMBOL(vmtruncate);

    /*
    * Primitive swap readahead code. We simply read an aligned block of
1630 * (1 << page_cluster) entries in the swap area. This method is chosen
    * because it doesn't cost us any seek time. We also make sure to queue
    * the 'original' request together with the readahead ones...
    *
    * This has been extended to use the NUMA policies from the mm triggering
1635 * the readahead.
    *
    * Caller must hold down_read on the vma->vm_mm if vma is not NULL.
    */

```

```

void swapin_readahead(swp_entry_t entry, unsigned long addr, struct vm_area_struct *vma)
1640 {
    #ifdef CONFIG_NUMA
        struct vm_area_struct *next_vma = vma ? vma->vm_next : NULL;
    #endif

    int i, num;
1645    struct page *new_page;
    unsigned long offset;

    /*
     * Get the number of handles we should do readahead io to.
     */
1650    num = valid_swaphandles(entry, &offset);
    for (i = 0; i < num; offset++, i++) {
        /* Ok, do the async read-ahead now */
        new_page = read_swap_cache_async(swp_entry(swp_type(entry),
1655                                     offset), vma, addr);

        if (!new_page)
            break;
        page_cache_release(new_page);
    #ifdef CONFIG_NUMA
1660        /*
         * Find the next applicable VMA for the NUMA policy.
         */
        addr += PAGE_SIZE;
        if (addr == 0)
1665            vma = NULL;
        if (vma) {
            if (addr >= vma->vm_end) {
                vma = next_vma;
                next_vma = vma ? vma->vm_next : NULL;
            }
            if (vma && addr < vma->vm_start)
                vma = NULL;
        } else {
1675            if (next_vma && addr >= next_vma->vm_start) {
                vma = next_vma;
                next_vma = vma->vm_next;
            }
        }
    #endif
1680    }
    lru_add_drain();
    /* Push any new pages onto the LRU now */
}

/*
1685 * We hold the mm semaphore and the page_table_lock on entry and
 * should release the pagetable lock on exit..
 */
static int do_swap_page(struct mm_struct * mm,
    struct vm_area_struct * vma, unsigned long address,
1690    pte_t *page_table, pmd_t *pmd, pte_t orig_pte, int write_access)
{
    struct page *page;
    swp_entry_t entry = pte_to_swp_entry(orig_pte);
    pte_t pte;
1695    int ret = VM_FAULT_MINOR;

    pte_unmap(page_table);
    spin_unlock(&mm->page_table_lock);
    page = lookup_swap_cache(entry);
1700    if (!page) {
        swapin_readahead(entry, address, vma);
    }
}

```

```

page = read_swap_cache_async(entry, vma, address);
if (!page) {
    /*
    * Back out if somebody else faulted in this pte while
    * we released the page table lock.
    */
    spin_lock(&mm->page_table_lock);
    page_table = pte_offset_map(pmd, address);
    if (likely(pte_same(*page_table, orig_pte)))
        ret = VM_FAULT_OOM;
    else
        ret = VM_FAULT_MINOR;
    pte_unmap(page_table);
    spin_unlock(&mm->page_table_lock);
    goto out;
}

/* Had to read the page from swap area: Major fault */
ret = VM_FAULT_MAJOR;
inc_page_state(pgmaifault);
grab_swap_token();
}

mark_page_accessed(page);
lock_page(page);

/*
* Back out if somebody else faulted in this pte while we
* released the page table lock.
*/
spin_lock(&mm->page_table_lock);
page_table = pte_offset_map(pmd, address);
if (unlikely(!pte_same(*page_table, orig_pte))) {
    pte_unmap(page_table);
    spin_unlock(&mm->page_table_lock);
    unlock_page(page);
    page_cache_release(page);
    ret = VM_FAULT_MINOR;
    goto out;
}

/* The page isn't present yet, go ahead with the fault. */

swap_free(entry);
if (vm_swap_full())
    remove_exclusive_swap_page(page);

mm->rss++;
acct_update_integrals();
update_mem_hiwater();

pte = mk_pte(page, vma->vm_page_prot);
if (write_access && can_share_swap_page(page)) {
    pte = maybe_mkdirty(pte_mkdirty(pte), vma);
    write_access = 0;
}
unlock_page(page);

flush_icache_page(vma, page);
set_pte(page_table, pte);
page_add_anon_rmap(page, vma, address);

if (write_access) {

```

```

1765         if (do_wp_page(mm, vma, address,
                        page_table, pmd, pte) == VM_FAULT_OOM)
            ret = VM_FAULT_OOM;
        goto out;
    }

1770     /* No need to invalidate – it was non-present before */
    update_mmu_cache(vma, address, pte);
    pte_unmap(page_table);
    spin_unlock(&mm->page_table_lock);

1775 out:
    return ret;
}

/*
1780 * We are called with the MM semaphore and page_table_lock
* spinlock held to protect against concurrent faults in
* multithreaded programs.
*/
static int
1785 do_anonymous_page(struct mm_struct *mm, struct vm_area_struct *vma,
                    pte_t *page_table, pmd_t *pmd, int write_access,
                    unsigned long addr)
{
    pte_t entry;
1790    struct page * page = ZERO_PAGE(addr);

    /* Read-only mapping of ZERO_PAGE. */
    entry = pte_wrprotect(mk_pte(ZERO_PAGE(addr), vma->vm_page_prot));

1795    /* ..except if it's a write access */
    if (write_access) {
        /* Allocate our own private page. */
        pte_unmap(page_table);
        spin_unlock(&mm->page_table_lock);

1800        if (unlikely(anon_vma_prepare(vma)))
            goto no_mem;
        page = alloc_zeroed_user_highpage(vma, addr);
        if (!page)
1805            goto no_mem;

        spin_lock(&mm->page_table_lock);
        page_table = pte_offset_map(pmd, addr);

1810        if (!pte_none(*page_table)) {
            pte_unmap(page_table);
            page_cache_release(page);
            spin_unlock(&mm->page_table_lock);
            goto out;
1815        }
        mm->rss++;
        acct_update_integrals();
        update_mem_hiwater();
        entry = maybe_mkdirty(mk_pte(page,
1820                                vma->vm_page_prot),
                                vma);

        lru_cache_add_active(page);
        SetPageReferenced(page);
        page_add_anon_rmap(page, vma, addr);

1825    }

    set_pte(page_table, entry);

```

```
pte_unmap(page_table);
```

```
1830 /* No need to invalidate – it was non-present before */  
update_mmu_cache(vma, addr, entry);  
spin_unlock(&mm->page_table_lock);
```

```
out:
```

```
return VM_FAULT_MINOR;
```

```
1835 no_mem:  
return VM_FAULT_OOM;
```

```
}
```

```
/*
```

```
1840 * do_no_page() tries to create a new page mapping. It aggressively  
* tries to share with existing pages, but makes a separate copy if  
* the "write_access" parameter is true in order to avoid the next  
* page fault.  
*
```

```
1845 * As this is called only for pages that do not currently exist, we  
* do not need to flush old virtual caches or the TLB.  
*
```

```
* This is called with the MM semaphore held and the page table  
* spinlock held. Exit with the spinlock released.
```

```
1850 */
```

```
static int
```

```
do_no_page(struct mm_struct *mm, struct vm_area_struct *vma,  
            unsigned long address, int write_access, pte_t *page_table, pmd_t *pmd)
```

```
{
```

```
1855 struct page * new_page;  
struct address_space *mapping = NULL;  
pte_t entry;  
unsigned int sequence = 0;  
int ret = VM_FAULT_MINOR;  
1860 int anon = 0;
```

```
if (!vma->vm_ops || !vma->vm_ops->nopage)  
    return do_anonymous_page(mm, vma, page_table,  
                             pmd, write_access, address);
```

```
1865 pte_unmap(page_table);  
spin_unlock(&mm->page_table_lock);
```

```
if (vma->vm_file) {  
    mapping = vma->vm_file->f_mapping;  
    sequence = mapping->truncate_count;  
    smp_rmb(); /* serializes i_size against truncate_count */  
}
```

```
retry:
```

```
cond_resched();  
1875 new_page = vma->vm_ops->nopage(vma, address & PAGE_MASK, &ret);
```

```
/*  
* No smp_rmb is needed here as long as there's a full  
* spin_lock/unlock sequence inside the ->nopage callback  
* (for the pagecache lookup) that acts as an implicit  
1880 * smp_mb() and prevents the i_size read to happen  
* after the next truncate_count read.  
*/
```

```
/* no page was available -- either SIGBUS or OOM */
```

```
1885 if (new_page == NOPAGE_SIGBUS)  
    return VM_FAULT_SIGBUS;
```

```
if (new_page == NOPAGE_OOM)  
    return VM_FAULT_OOM;
```

```
1890 /*
```

```

1895     * Should we do an early C-O-W break?
1896     */
1897     if (write_access && !(vma->vm_flags & VM_SHARED)) {
1898         struct page *page;
1899
1900         if (unlikely(anon_vma_prepare(vma)))
1901             goto oom;
1902         page = alloc_page_vma(GFP_HIGHUSER, vma, address);
1903         if (!page)
1904             goto oom;
1905         copy_user_highpage(page, new_page, address);
1906         page_cache_release(new_page);
1907         new_page = page;
1908         anon = 1;
1909     }
1910
1911     spin_lock(&mm->page_table_lock);
1912     /*
1913      * For a file-backed vma, someone could have truncated or otherwise
1914      * invalidated this page. If unmap_mapping_range got called,
1915      * retry getting the page.
1916      */
1917     if (mapping && unlikely(sequence != mapping->truncate_count)) {
1918         sequence = mapping->truncate_count;
1919         spin_unlock(&mm->page_table_lock);
1920         page_cache_release(new_page);
1921         goto retry;
1922     }
1923     page_table = pte_offset_map(pmd, address);
1924
1925     /*
1926      * This silly early PAGE_DIRTY setting removes a race
1927      * due to the bad i386 page protection. But it's valid
1928      * for other architectures too.
1929      */
1930     /* Note that if write_access is true, we either now have
1931      * an exclusive copy of the page, or this is a shared mapping,
1932      * so we can make it writable and dirty to avoid having to
1933      * handle that later.
1934      */
1935     /* Only go through if we didn't race with anybody else... */
1936     if (pte_none(*page_table)) {
1937         if (!PageReserved(new_page))
1938             ++mm->rss;
1939         acct_update_integrals();
1940         update_mem_hiwater();
1941
1942         flush_icache_page(vma, new_page);
1943         entry = mk_pte(new_page, vma->vm_page_prot);
1944         if (write_access)
1945             entry = maybe_mkwritable(pte_mkdirty(entry), vma);
1946         set_pte(page_table, entry);
1947         if (anon) {
1948             lru_cache_add_active(new_page);
1949             page_add_anon_rmap(new_page, vma, address);
1950         } else
1951             page_add_file_rmap(new_page);
1952         pte_unmap(page_table);
1953     } else {
1954         /* One of our sibling threads was faster, back out. */
1955         pte_unmap(page_table);
1956         page_cache_release(new_page);
1957         spin_unlock(&mm->page_table_lock);

```

```

1955         goto out;
    }

    /* no need to invalidate: a not-present page shouldn't be cached */
    update_mmu_cache(vma, address, entry);
    spin_unlock(&mm->page_table_lock);

1960 out:
    return ret;

    oom:
    page_cache_release(new_page);
    ret = VM_FAULT_OOM;
1965     goto out;
}

/*
 * Fault of a previously existing named mapping. Repopulate the pte
1970 * from the encoded file_pte if possible. This enables swappable
 * nonlinear vmas.
 */
static int do_file_page(struct mm_struct * mm, struct vm_area_struct * vma,
    unsigned long address, int write_access, pte_t *pte, pmd_t *pmd)
1975 {
    unsigned long pgoff;
    int err;

    BUG_ON(!vma->vm_ops || !vma->vm_ops->nopage);
1980 /*
 * Fall back to the linear mapping if the fs does not support
 * ->populate:
 */
    if (!vma->vm_ops || !vma->vm_ops->populate ||
1985         (write_access && !(vma->vm_flags & VM_SHARED))) {
        pte_clear(pte);
        return do_no_page(mm, vma, address, write_access, pte, pmd);
    }

1990     pgoff = pte_to_pgoff(*pte);

    pte_unmap(pte);
    spin_unlock(&mm->page_table_lock);

1995     err = vma->vm_ops->populate(vma, address & PAGE_MASK, PAGE_SIZE, vma->vm_page_prot, pgoff,
    ↪ 0);
    if (err == -ENOMEM)
        return VM_FAULT_OOM;
    if (err)
        return VM_FAULT_SIGBUS;
2000     return VM_FAULT_MAJOR;
}

/*
 * These routines also need to handle stuff like marking pages dirty
2005 * and/or accessed for architectures that don't do it in hardware (most
 * RISC architectures). The early dirtying is also good on the i386.
 *
 * There is also a hook called "update_mmu_cache()" that architectures
 * with external mmu caches can use to update those (ie the Sparc or
2010 * PowerPC hashed page tables that act as extended TLBs).
 *
 * Note the "page_table_lock". It is to protect against kswapd removing
 * pages from under us. Note that kswapd only ever __removes__ pages, never
 * adds them. As such, once we have noticed that the page is not present,
2015 * we can drop the lock early.

```



```

*
* The adding of pages is protected by the MM semaphore (which we hold),
* so we don't need to worry about a page being suddenly been added into
* our VM.
2020 *
* We enter with the pagetable spinlock held, we are supposed to
* release it when done.
*/
static inline int handle_pte_fault(struct mm_struct *mm,
2025 struct vm_area_struct * vma, unsigned long address,
int write_access, pte_t *pte, pmd_t *pmd)
{
    pte_t entry;

2030 entry = *pte;
    if (!pte_present(entry)) {
        /*
         * If it truly wasn't present, we know that kswapd
         * and the PTE updates will not touch it later. So
2035         * drop the lock.
        */
        if (pte_none(entry))
            return do_no_page(mm, vma, address, write_access, pte, pmd);
        if (pte_file(entry))
2040            return do_file_page(mm, vma, address, write_access, pte, pmd);
        return do_swap_page(mm, vma, address, pte, pmd, entry, write_access);
    }

    if (write_access) {
2045         if (!pte_write(entry))
            return do_wp_page(mm, vma, address, pte, pmd, entry);

        entry = pte_mkdirty(entry);
    }
2050 entry = pte_mkyoung(entry);
    ptep_set_access_flags(vma, address, pte, entry, write_access);
    update_mmu_cache(vma, address, entry);
    pte_unmap(pte);
    spin_unlock(&mm->page_table_lock);
2055 return VM_FAULT_MINOR;
}

/*
* By the time we get here, we already hold the mm semaphore
2060 */
int handle_mm_fault(struct mm_struct *mm, struct vm_area_struct * vma,
unsigned long address, int write_access)
{
    pgd_t *pgd;
2065 pud_t *pud;
    pmd_t *pmd;
    pte_t *pte;

    __set_current_state(TASK_RUNNING);

2070 inc_page_state(pgfault);

    if (is_vm_hugetlb_page(vma))
        return VM_FAULT_SIGBUS; /* mapping truncation does this. */
2075
    /*
     * We need the page table lock to synchronize with kswapd
     * and the SMP-safe atomic PTE updates.

```

```

2080     */
    pgd = pgd_offset(mm, address);
    spin_lock(&mm->page_table_lock);

    pud = pud_alloc(mm, pgd, address);
    if (!pud)
2085         goto oom;

    pmd = pmd_alloc(mm, pud, address);
    if (!pmd)
        goto oom;

2090    pte = pte_alloc_map(mm, pmd, address);
    if (!pte)
        goto oom;

2095    return handle_pte_fault(mm, vma, address, write_access, pte, pmd);

oom:
    spin_unlock(&mm->page_table_lock);
    return VM_FAULT_OOM;
2100 }

#ifdef __ARCH_HAS_4LEVEL_HACK
/*
 * Allocate page upper directory.
2105 *
 * We've already handled the fast-path in-line, and we own the
 * page table lock.
 *
 * On a two-level or three-level page table, this ends up actually being
2110 * entirely optimized away.
 */
pud_t fastcall * __pud_alloc(struct mm_struct *mm, pgd_t *pgd, unsigned long address)
{
    pud_t *new;

2115    spin_unlock(&mm->page_table_lock);
    new = pud_alloc_one(mm, address);
    spin_lock(&mm->page_table_lock);
    if (!new)
2120         return NULL;

    /*
     * Because we dropped the lock, we should re-check the
     * entry, as somebody else could have populated it..
2125 */
    if (pgd_present(*pgd)) {
        pud_free(new);
        goto out;
    }
    pgd_populate(mm, pgd, new);
2130 out:
    return pud_offset(pgd, address);
}

2135 /*
 * Allocate page middle directory.
 *
 * We've already handled the fast-path in-line, and we own the
 * page table lock.
2140 *
 * On a two-level page table, this ends up actually being entirely

```

```

    * optimized away.
    */
pmd_t fastcall * __pmd_alloc(struct mm_struct * mm, pud_t * pud, unsigned long address)
2145 {
    pmd_t * new;

    spin_unlock(&mm->page_table_lock);
    new = pmd_alloc_one(mm, address);
2150 spin_lock(&mm->page_table_lock);
    if (!new)
        return NULL;

    /*
2155  * Because we dropped the lock, we should re-check the
    * entry, as somebody else could have populated it..
    */
    if (pud_present(* pud)) {
        pmd_free(new);
2160 goto out;
    }
    pud_populate(mm, pud, new);
out:
    return pmd_offset(pud, address);
2165 }
#else
pmd_t fastcall * __pmd_alloc(struct mm_struct * mm, pud_t * pud, unsigned long address)
{
    pmd_t * new;
2170
    spin_unlock(&mm->page_table_lock);
    new = pmd_alloc_one(mm, address);
    spin_lock(&mm->page_table_lock);
    if (!new)
2175 return NULL;

    /*
    * Because we dropped the lock, we should re-check the
    * entry, as somebody else could have populated it..
2180  */
    if (pgd_present(* pud)) {
        pmd_free(new);
        goto out;
    }
2185 pgd_populate(mm, pud, new);
out:
    return pmd_offset(pud, address);
}
#endif
2190
int make_pages_present(unsigned long addr, unsigned long end)
{
    int ret, len, write;
    struct vm_area_struct * vma;
2195
    vma = find_vma(current->mm, addr);
    if (!vma)
        return -1;
    write = (vma->vm_flags & VM_WRITE) != 0;
2200 if (addr >= end)
    BUG();
    if (end > vma->vm_end)
    BUG();
    len = (end+PAGE_SIZE-1)/PAGE_SIZE-addr/PAGE_SIZE;

```

```

2205         ret = get_user_pages(current, current->mm, addr,
                                len, write, 0, NULL, NULL);
        if (ret < 0)
            return ret;
        return ret == len ? 0 : -1;
2210     }

    /*
     * Map a vmalloc()-space virtual address to the physical page.
     */
2215     struct page * vmalloc_to_page(void * vmalloc_addr)
    {
        unsigned long addr = (unsigned long) vmalloc_addr;
        struct page * page = NULL;
        pgd_t * pgd = pgd_offset_k(addr);
2220         pud_t * pud;
        pmd_t * pmd;
        pte_t * ptep, pte;

        if (!pgd_none(*pgd)) {
2225             pud = pud_offset(pgd, addr);
            if (!pud_none(*pud)) {
                pmd = pmd_offset(pud, addr);
                if (!pmd_none(*pmd)) {
2230                     ptep = pte_offset_map(pmd, addr);
                    pte = *ptep;
                    if (pte_present(pte))
                        page = pte_page(pte);
                    pte_unmap(ptep);
                }
2235             }
        }
        return page;
    }

2240 EXPORT_SYMBOL(vmalloc_to_page);

    /*
     * Map a vmalloc()-space virtual address to the physical page frame number.
     */
2245     unsigned long vmalloc_to_pfn(void * vmalloc_addr)
    {
        return page_to_pfn(vmalloc_to_page(vmalloc_addr));
    }

2250 EXPORT_SYMBOL(vmalloc_to_pfn);

    /*
     * update_mem_hiwater
     *   - update per process rss and vm high water data
2255     */
    void update_mem_hiwater(void)
    {
        struct task_struct * tsk = current;

2260         if (tsk->mm) {
            if (tsk->mm->hiwater_rss < tsk->mm->rss)
                tsk->mm->hiwater_rss = tsk->mm->rss;
            if (tsk->mm->hiwater_vm < tsk->mm->total_vm)
                tsk->mm->hiwater_vm = tsk->mm->total_vm;
2265         }
    }

```

```

#if !defined(__HAVE_ARCH_GATE_AREA)

2270 #if defined(AT_SYSINFO_EHDR)
    struct vm_area_struct gate_vma;

    static int __init gate_vma_init(void)
    {
2275         gate_vma.vm_mm = NULL;
        gate_vma.vm_start = FIXADDR_USER_START;
        gate_vma.vm_end = FIXADDR_USER_END;
        gate_vma.vm_page_prot = PAGE_READONLY;
        gate_vma.vm_flags = 0;
2280         return 0;
    }
    __initcall(gate_vma_init);
#endif

2285 struct vm_area_struct *get_gate_vma(struct task_struct *tsk)
    {
        #ifdef AT_SYSINFO_EHDR
            return &gate_vma;
        #else
2290         return NULL;
        #endif
    }

    int in_gate_area_no_task(unsigned long addr)
2295 {
        #ifdef AT_SYSINFO_EHDR
            if ((addr >= FIXADDR_USER_START) && (addr < FIXADDR_USER_END))
                return 1;
        #endif
2300         return 0;
    }

2303 #endif /* __HAVE_ARCH_GATE_AREA */

```